

Trace-based capture and replay for detecting Go concurrency bugs

Study project

Erik Kassubek


Department of Computer Science
Albert Ludwig University of Freiburg

04.06.2025

Advisor: Prof. Dr. Sulzmann
Examiner: Prof. Dr. Thiemann


Concurrency Bugs

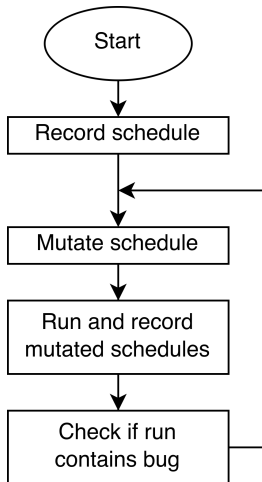
```
1 // mixed deadlock
2 // simplified from Kubernetes
3 var (
4     m = sync.Mutex{}
5     c = make(chan bool)
6 )
7
8 func R1() {
9     for i := 0; i < 2; i++ {
10         <-c
11         m.Lock()
12         m.Unlock()
13     }
14 }
15
16 func R2() {
17     m.Lock()
18     c <- true
19     m.Unlock()
20 }
21
22 func main() {
23     go R1() // G1
24     go R2() // G2
25     R2()   // G3
26 }
```



Concurrency Bugs

```
1 // mixed deadlock
2 // simplified from Kubernetes
3 var (
4     m = sync.Mutex{}
5     c = make(chan bool)
6 )
7
8 func R1() {
9     for i := 0; i < 2; i++ {
10         <-c
11         m.Lock()
12         m.Unlock()
13     }
14 }
15
16 func R2() {
17     m.Lock()
18     c <- true
19     m.Unlock()
20 }
21
22 func main() {
23     go R1() // G1
24     go R2() // G2
25     R2()   // G3
26 }
```





Project

- Implement/improve execution tracing of go programs
- Implement replay mechanics for go programs
- Implement simple fuzzing approaches

📁 advocate	redraw gfuzz mut	1 hour ago
📁 doc	simplify trace format for select	2 days ago
📁 doc_proj	simplify trace format for select	2 days ago
📁 examples	add achievements to Project	4 days ago
📁 go-patch	temp fix for stuck trace write	13 hours ago
📄 .gitignore	remove eval folder, update go-patch	last month
📄 LICENSE	first commit	2 years ago
📄 PROJECT.md	add pre	4 days ago
📄 README.md	improve simple leak detection	3 weeks ago

Project

- Implement/improve execution tracing of go programs
- Implement replay mechanics for go programs
- Implement simple fuzzing approaches

📁 advocate	redraw gfuzz mut	1 hour ago
📁 doc	simplify trace format for select	2 days ago
📁 doc_proj	simplify trace format for select	2 days ago
📁 examples	add achievements to Project	4 days ago
📁 go-patch	temp fix for stuck trace write	13 hours ago
📄 .gitignore	remove eval folder, update go-patch	last month
📄 LICENSE	first commit	2 years ago
📄 PROJECT.md	add pre	4 days ago
📄 README.md	improve simple leak detection	3 weeks ago

Project

- Implement/improve execution tracing of go programs
- Implement replay mechanics for go programs
- Implement simple fuzzing approaches

📁 advocate	redraw gfuzz mut	1 hour ago
📁 doc	simplify trace format for select	2 days ago
📁 doc_proj	simplify trace format for select	2 days ago
📁 examples	add achievements to Project	4 days ago
📁 go-patch	temp fix for stuck trace write	13 hours ago
📄 .gitignore	remove eval folder, update go-patch	last month
📄 LICENSE	first commit	2 years ago
📄 PROJECT.md	add pre	4 days ago
📄 README.md	improve simple leak detection	3 weeks ago

Tracing

- Record concurrency operations
 - fork, routine end
 - mutex
 - channel, select
 - once, wait group, conditional variables
 - atomic operations
- Implemented in runtime
- Routine local recording

- When was operation executed
- Required to reconstruct global trace
- Implemented as `atomic.Int32`

Tracing

Elements

```
1 func (m *Mutex) Lock() { // simplified
2     advocateIndex := runtime.AdvocateMutexPre(m.id, runtime.OperationMutexLock)
3
4     m.mu.Lock()
5
6     runtime.AdvocateMutexPost(advocateIndex, true)
7 }
```

```
1 type g struct { // simplified
2     ... // routine data, e.g. stack parameter
3
4     advocateRoutineInfo *AdvocateRoutine
5 }
```

Replay

- Execute program schedule based on trace
- Wait and release mechanism
- Keep track of next operation to be executed
- Operations wait until it is their turn

Replay

Key

```
1 func BuildReplayKey(routine int, file string, line int) string {  
2     return intToString(routine) + ":" + file + ":" + intToString(line)  
3 }
```

```
1 c := make(chan int, 2)  
2  
3 for i := range 2 {  
4     c <- 1  
5 }  
6  
7 a := func() {  
8     <- c  
9 }  
10  
11 go a()  
12 go a()
```

- Note: Inlining has been deactivated

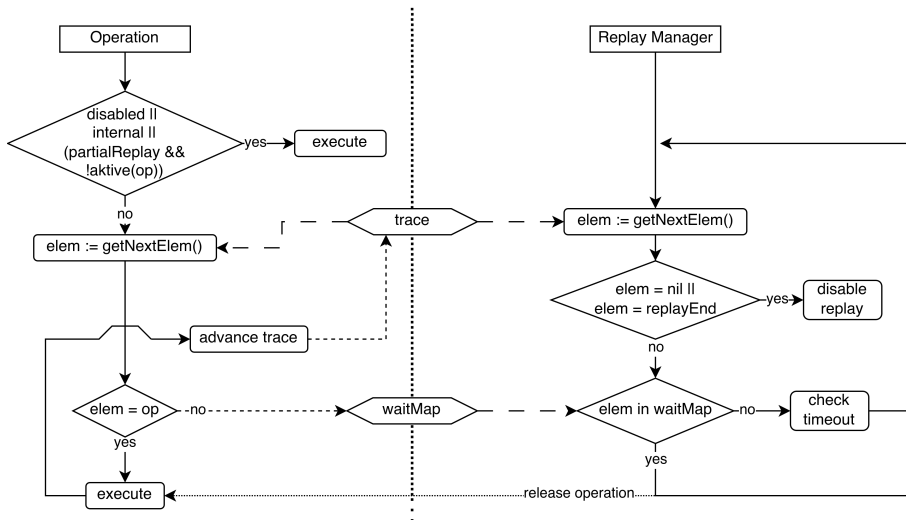
Replay

In Operation

```
1 func (m *Mutex) Lock() { // simplified
2     active, chWait, chAck := runtime.WaitForReplay(2, true)
3
4     if active {
5         defer func() { chAck <- struct{}{} }()
6         replayElem := <-chWait
7     }
8
9     ...
10
```

Replay

Replay



Replay

Select

```
1 select {
2   case c <- 1:
3   case <- d:
4   case <- e:
5   case f <- 1:
6 }
```

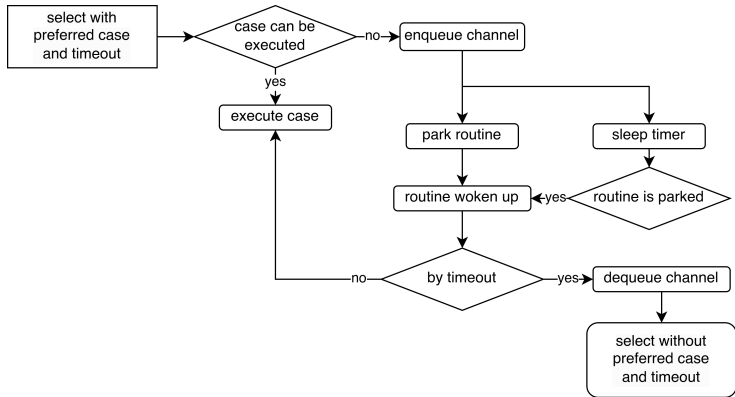
```
1 scases = [
2   {c, send},
3   {f, send},
4   {e, recv},
5   {d, recv}
6 ]
7
8 selectedIndex = 2
```

```
1 // simplified
2 for _, casi := range pollorder {
3   cas = &scases[casi] // *scase
4   c = cas.c           // *hchan
5
6   if casi != selectedIndex {
7     continue
8   }
9
10  if casi >= nsends {
11    ... // enqueue as recv
12  } else {
13    ... // enqueue as send
14  }
15
16  ...
17 }
```

Replay

Select with timeout

- First try only preferred case
- On timeout, if no partner is available, try all
- Required for some fuzzing approaches



Fuzzing

- Based on work by Lio et al. [1]
- Reorder concurrent communication with selects

¹Z. Liu et al., "Who goes first? detecting go concurrency bugs via message reordering"

Fuzzing

GFuzz: Example

```
1 func main() {
2     c := make(chan int, 2)
3     d := make(chan int, 2)
4     e := make(chan int, 1)
5
6     go func() {
7         c <- 1
8         c <- 1
9     }()
10
11    go func() {
12        d <- 1
13        d <- 1
14    }()
15
16    select {
17        case <-c:
18        case <-d:
19            close(e)
20    }
21
22    select {
23        case <-c:
24        case <-d:
25            e <- 1
26    }
27 }
```

- For each select, get executed case
- Mutate preferred cases
- Replay select with timeout
- No replay of operation order

- Based on work by Jiang et al. [2]
- Improvement of GFuzz
- Mutate order of channel and mutex
- Implement original and improved version

²Z. Jiang et al., “Effective Concurrency Testing for Go via Directional Primitive-Constrained Interleaving Exploration”

- Mutate order of concurrency operations
- Construct relations between operations
- Rule based mutation on subset (scheduling chains) of operations
 - Abridge
 - Flip
 - Substitute
 - Augment
- Partial replay of subset

Fuzzing

GoPie+: Improvements I

```
1 func main() {
2     c := make(chan bool)
3
4     go func() {
5         c <- true
6     }
7
8     go func() {
9         c <- false
10    }
11
12    if res := <- c; res {
13        // code that contains mutated section
14        ...
15    }
16 }
```

- full replay before mutated operations to guaranty that the program reaches the modified code block
- full coverage of all concurrency primitives
- consider scheduling chains resulting from mutated program runs
- mutate partially executed operations
- exclude superfluous mutations

Fuzzing

GoPie: GoBench

- Apply GoPie to GoBench [3]
- Number of found bugs:

Paper	GoPie	GoPie+
66	66	69

³T. Yuan et al., “GoBench: A Benchmark Suite of Real-World Go Concurrency Bugs”

Summary

- Recording
- Replay
- Fuzzing
 - GFuzz
 - GoPie / GoPie+

- [1] Z. Liu, S. Xia, Y. Liang, L. Song, and H. Hu, "Who goes first? detecting go concurrency bugs via message reordering," in *Proceedings of the 27th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*, ser. ASPLOS '22, Lausanne, Switzerland: Association for Computing Machinery, 2022, pp. 888–902, ISBN: 9781450392051. DOI: 10.1145/3503222.3507753. [Online]. Available: <https://doi.org/10.1145/3503222.3507753>.
- [2] Z. Jiang, M. Wen, Y. Yang, C. Peng, P. Yang, and H. Jin, "Effective concurrency testing for go via directional primitive-constrained interleaving exploration," in *Proceedings of the 38th IEEE/ACM International Conference on Automated Software Engineering*, ser. ASE '23, Echternach, Luxembourg: IEEE Press, 2024, pp. 1364–1376, ISBN: 9798350329964. DOI: 10.1109/ASE56229.2023.00086. [Online]. Available: <https://doi.org/10.1109/ASE56229.2023.00086>.
- [3] T. Yuan, G. Li, J. Lu, C. Liu, L. Li, and J. Xue, "Gobench: A benchmark suite of real-world go concurrency bugs," in *2021 IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*, 2021, pp. 187–199. DOI: 10.1109/CGO51591.2021.9370317.

Appendix

Atomic Old

```
1 // doc.go
2 func AddInt32(addr *int32, delta int32) (new int32)
3
4 // type.go
5 func (x *Int32) Add(delta int32) (new int32) {
6     return AddInt32(&x.v, delta)
7 }
8
9 // asm.s
10 TEXT ·AddInt32(SB),NOSPLIT,$0
11     JMP internal/runtime/atomic·Xadd(SB)
```

Appendix

Atomic New

```
1 // advocate_atomic.go
2 func AddInt32(addr *int32, delta int32) (new int32) { // simplified
3     runtime.AdvocateAtomic(addr, runtime.OperationAtomicAdd, 2)
4     return AddInt32Advocate(addr, delta)
5 }
6
7 // type.go
8 func (x *Int32) Add(delta int32) (new int32) {
9     return AddInt32AdvocateType(&x.v, delta)
10 }
11
12 // advocate_atomic_type.go
13 func AddInt32AdvocateType(addr *int32, delta int32) (new int32) { // simplified
14     runtime.AdvocateAtomic(addr, runtime.OperationAtomicAdd, 3)
15     return AddInt32Advocate(addr, delta)
16 }
17
18 // doc.go
19 func AddInt32Advocate(addr *int32, delta int32) (new int32)
20
21 // asm.s
22 TEXT ·AddInt32Advocate(SB),NOSPLIT,$0
23     JMP internal/runtime/atomic·Xadd(SB)
```

Tracing: Routine Local Tracing

```
1 type g struct { // simplified
2     ... // routine data, e.g. stack parameter
3
4     advocateRoutineInfo *AdvocateRoutine
5 }
```

```
1 type AdvocateRoutine struct {
2     id          uint64
3     Trace       []traceElement
4     G           *g
5     replayID    int
6     maxObjectId uint64
7 }
```


Appendix

Tracing: Elements

```
1  type Mutex struct {
2      _ noCopy
3
4      mu isync.Mutex
5
6      id uint64 // id for the recording
7  }
8
9  func (m *Mutex) Lock() { // simplified
10     if m.id == 0 {
11         m.id = runtime.GetAdvocateObjectID()
12     }
13
14     advocateIndex := runtime.AdvocateMutexPre(m.id, runtime.OperationMutexLock)
15
16     m.mu.Lock()
17
18     runtime.AdvocateMutexPost(advocateIndex, true)
19 }
```

Appendix

Replay: Replay Manager

Algorithm 1: ReplayManager

```
1 while true do
2   if waitForAck  $\neq$  nil then
3     |   waitAck();
4     |   nextElem();
5   nElem  $\leftarrow$  nextElemInTrace();
6   if nElem == nil || nElem == ReplayEnd then
7     |   return
8   if elem, ok := waitMap[nElem.key]; ok then
9     |   elem.chWait.send(nextElem);
10    |   if elem.waitForAck then
11      |   |   waitAck();
12      |   |   nextElem();
13    if len(waitMap) > 0 && timeSinceLastRelease > timeoutTime then
14      |   timeout();
```

Appendix: Replay: WaitForReplay

Algorithm 2: WaitForReplay(*op*)

```
1 if replayDisabled // ignored then
2   |   return false, nil, nil;
3 chWait ← make(chan ReplayElem, 1);
4 chAck ← make(chan struct, 1);
5 elem ← nextElemInTrace();
6 if match(op, elem) then
7   |   chWait.send(elem);
8   |   if op.WaitForAck then
9     |   waitForAck ← {op, chAck};
10  |   else
11    |   nextElem();
12 else
13   |   waitMap[op.key] ← {op, chWait, chAck};
14 return true, chWait, chAck;
```

Replay

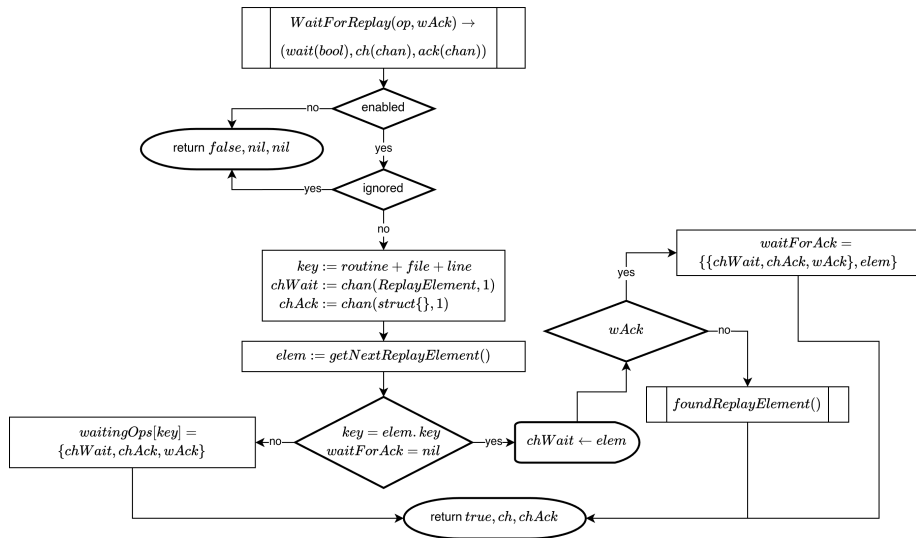
Key

```
1 func BuildReplayKey(routine int, file string, line int) string {  
2     return intToString(routine) + ":" + file + ":" + intToString(line)  
3 }
```

```
1 c := make(chan int, 2)  
2  
3 for i := range 2 {  
4     c <- 1  
5 }  
6  
7 a := func() {  
8     <- c  
9 }  
10  
11 go a()  
12 go a()
```

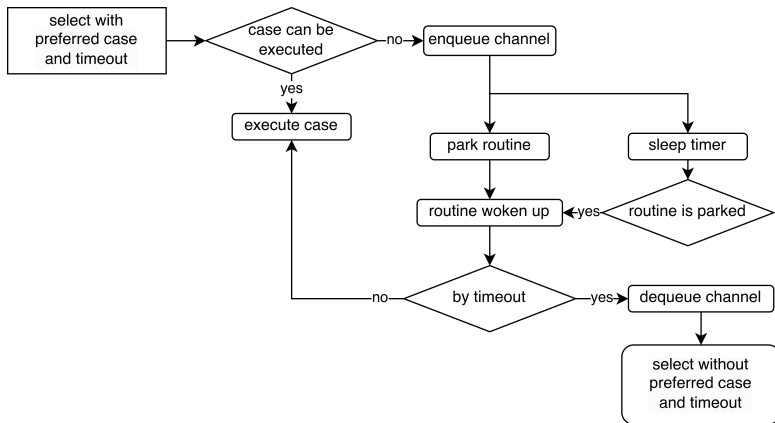
- Note: Inlining has been deactivated

Appendix: Replay: In Operation II



Replay

Select with timeout



Appendix

Select Timeout I

```
1 func selectgo(...) (int, bool) { // simplified
2     fuzzingEnabled, fuzzingIndex, timeout := AdvocateFuzzingGetPreferredCase(2)
3
4     if fuzzingEnabled {
5         if ok, i, b, _ := selectWithPrefCase(..., timeout); ok {
6             return i, b
7         }
8     }
9
10    return originalSelect(...)
11 }
```

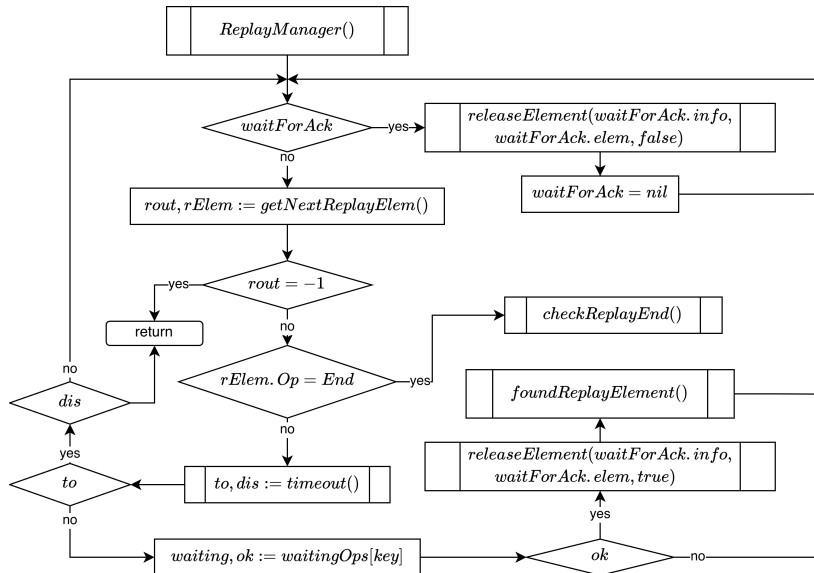
Appendix

Select Timeout II

```
1 func goparkWithTimeout(..., timeout int64) { // simplified
2     mp := acquirem()
3     gp := mp.curg
4
5     // Setup timer if timeout is non-zero
6     if timeout > 0 {
7         go func() {
8             sleep(float64(timeout))
9             if readgstatus(gp) == _Gwaiting {
10                 goready(gp, traceskip)
11             }
12         }()
13     }
14
15     // Run the original gopark logic
16     ...
```


Appendix

Replay: Replay Manager



Appendix

Replay: Release

```
1 // simplified
2 func releaseElement(elem replayChan, elemReplay ReplayElement, rel bool) bool {
3     if rel {
4         elem.chWait <- elemReplay
5         elem.released = true
6     }
7
8     if elem.waitAck {
9         select {
10             case <-elem.chAck:
11             case <-after(acknowledgementMaxWait):
12                 ... // store info that timeout
13         }
14     }
15
16     ... // some other controls
17 }
```

Appendix

Replay: Timeout

- Program may get stuck
 - Something went wrong
 - E.g. error in trace rewrite, random element in program
- Only one element may be the problem
- Try to fix by
 - Release oldest waiting element
 - Skip one element in trace
- If it cannot be fixed
 - disable replay

Appendix

Fuzzing: Select with preferred case

```
1 func selectgo(...) (int, bool) { // simplified
2     fuzzingEnabled, fuzzingIndex, timeout := AdvocateFuzzingGetPreferredCase(2)
3
4     if fuzzingEnabled {
5         if ok, i, b, _ := selectWithPrefCase(..., timeout); ok {
6             return i, b
7         }
8     }
9
10    return originalSelect(...)
11 }
```

Appendix

Fuzzing: Gopark with timeout

```
1 func goparkWithTimeout(..., timeout int64) { // simplified
2     mp := acquirem()
3     gp := mp.curg
4
5     if timeout > 0 {
6         go func() {
7             sleep(float64(timeout))
8             if readgstatus(gp) == _Gwaiting {
9                 goready(gp, traceskip)
10            }
11        }()
12    }
13
14    ...
```

- Mutate order of concurrency operations
- Scheduling chain:


$$SC = \{\langle o_1, r_1 \rangle, \dots, \langle o_n, r_n \rangle \mid r_i \neq r_{i+1}, 1 \leq i \leq n - 1\}$$

- Relations:
 - Rule 1: $c' \in Rel_1(c)$ if c and c' neighbors in same routine
 - Rule 2: $c' \in Rel_2(c)$ if c and c' neighbors in different routines but on same primitive
 - Rule 3: $c' \in Rel_1(c), c'' \in Rel_2(c') \rightarrow c'' \in Rel_2(c)$
 - Rule 4: $c' \in Rel_2(c), c'' \in Rel_2(c') \rightarrow c'' \in Rel_2(c)$

- Mutate scheduling chains
 - Abridge
 - $\{o_i, o_{i+1}, \dots, o_{j-1}, o_j\} \Rightarrow \{o_{i+1}, \dots, o_{j-1}, o_j\}, \{o_i, o_{i+1}, \dots, o_{j-1}\}$
 - Flip (in paper, not in implementation)
 - $\{\dots, o_i, o_j, \dots\} \Rightarrow \{\dots, o_j, o_i, \dots\}$
 - Substitute
 - $o_j \in Rel_1(o_i), \{\dots, o_i, \dots\} \Rightarrow \{\dots, o_j, \dots\}$
 - Augment
 - $o_j \in Rel_2(o_i), \{\dots, o_i\} \Rightarrow \{\dots, o_i, o_j\}$

Appendix


Fuzzing: GoPie: Example



Step	Mutation	Scheduling chain
------	----------	------------------

Appendix


Fuzzing: GoPie: Example



Step	Mutation	Scheduling chain
Init	-	$\{\langle G2, 19 \rangle, \langle G1, 11 \rangle\}$

Appendix


Fuzzing: GoPie: Example



Step	Mutation	Scheduling chain
Init	-	$\{\langle G2, 19 \rangle, \langle G1, 11 \rangle\}$
Step 1	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 11 \rangle\}$

Appendix


Fuzzing: GoPie: Example



Step	Mutation	Scheduling chain
Init	-	$\{\langle G2, 19 \rangle, \langle G1, 11 \rangle\}$
Step 1	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 11 \rangle\}$
Step 2	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 12 \rangle\}$

Appendix


Fuzzing: GoPie: Example



Step	Mutation	Scheduling chain
Init	-	$\{\langle G2, 19 \rangle, \langle G1, 11 \rangle\}$
Step 1	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 11 \rangle\}$
Step 2	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 12 \rangle\}$
Step 3	Augment	$\{\langle G2, 18 \rangle, \langle G1, 12 \rangle, \langle G3, 17 \rangle\}$

Appendix


Fuzzing: GoPie: Example



Step	Mutation	Scheduling chain
Init	-	$\{\langle G2, 19 \rangle, \langle G1, 11 \rangle\}$
Step 1	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 11 \rangle\}$
Step 2	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 12 \rangle\}$
Step 3	Augment	$\{\langle G2, 18 \rangle, \langle G1, 12 \rangle, \langle G3, 17 \rangle\}$
Step 4	Flip	$\{\langle G2, 18 \rangle, \langle G3, 17 \rangle, \langle G1, 12 \rangle\}$

Appendix


Fuzzing: GoPie: Example



Step	Mutation	Scheduling chain
Init	-	$\{\langle G2, 19 \rangle, \langle G1, 11 \rangle\}$
Step 1	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 11 \rangle\}$
Step 2	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 12 \rangle\}$
Step 3	Augment	$\{\langle G2, 18 \rangle, \langle G1, 12 \rangle, \langle G3, 17 \rangle\}$
Step 4	Flip	$\{\langle G2, 18 \rangle, \langle G3, 17 \rangle, \langle G1, 12 \rangle\}$
Step 5	Substitute	$\{\langle G2, 18 \rangle, \langle G3, 17 \rangle, \langle G1, 11 \rangle\}$

Appendix

Fuzzing: GoPie: Example



Step	Mutation	Scheduling chain
Init	-	$\{\langle G2, 19 \rangle, \langle G1, 11 \rangle\}$
Step 1	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 11 \rangle\}$
Step 2	Substitute	$\{\langle G2, 18 \rangle, \langle G1, 12 \rangle\}$
Step 3	Augment	$\{\langle G2, 18 \rangle, \langle G1, 12 \rangle, \langle G3, 17 \rangle\}$
Step 4	Flip	$\{\langle G2, 18 \rangle, \langle G3, 17 \rangle, \langle G1, 12 \rangle\}$
Step 5	Substitute	$\{\langle G2, 18 \rangle, \langle G3, 17 \rangle, \langle G1, 11 \rangle\}$